

Cython: A Quick Overview

By Jason Paryani

Presentation will be available at:

github.com/jparyani/cython-presentation

Intro

- “**Cython** is a language that makes writing C extensions for the Python language as easy as **Python** itself.”
(www.cython.org)
- This makes it easy to:
 - Call directly into C libraries quickly and easily
 - Speed up slow parts of your code
 - Let's you mix Python and C, quickly and easily

Caveats

- If you aren't familiar with C/C++, parts of this presentation will go over your head.
- I'm going to make extensive use of micro-benchmarks
 - Only intended to give you a vague idea of the speed differences
 - gcc was used with `-O3`
 - Don't succumb to pre-mature optimization
 - Check out http://docs.cython.org/src/tutorial/profiling_tutorial.html
- Cython isn't for everything

Syntax

- Pretty similar to Python, with just a few extra things thrown in
- Types
- Structs/Unions/Enums
- Function definitions
- Extension Classes
- Extern blocks
- Lots of other stuff that we'll ignore for now

Types

```
cdef int i, j, k
```

```
cdef float f, g[42], *h
```

- cdef followed by a valid type signifies a type a variable defined with a type
- All the basic types from C are there ([unsigned] char, int, long, long long, float, double) and object
- You can typedef your own:

```
ctypedef unsigned long Ulong
```

- Also there's:

```
from libc.stdint cimport uint32_t
```

Casting Types

- Casting is done through the < > operator

```
from libc.stdlib cimport atoi
```

```
def parse_int(string):  
    return atoi(<char *>string)
```

- Python strings->char * are actually a special case
 - Read the docs

Structs/Unions/Enums

cdef struct Grail:

int age

float volume

cdef union Food:

char *spam

float *eggs

cdef enum CheeseType:

cheddar, edam, camembert

Functions

Python definition:

```
def eggs(unsigned long l, f): ...
```

Cython definition:

```
cdef int eggs(unsigned long l, float f): ...
```

Hybrid definition:

```
cpdef int eggs(unsigned long l, float f): ...
```

Equivalent to a union of the 2 above definitions

Function Call Overhead

Cython

```
cdef _test():  
    return 1  
def test(int n):  
    cdef int i  
    for i in range(n):  
        _test()
```

Compiled with -O0

timeit test(100)

1000000 loops, best of 3: **1.16 us** per loop

timeit test(1000)

100000 loops, best of 3: **9.83 us** per loop

Python

```
def _test():  
    return 1  
def test(n):  
    for i in range(n):  
        _test()
```

timeit test(100)

100000 loops, best of 3: **16.1 us** per loop

timeit test(1000)

10000 loops, best of 3: **155 us** per loop

Tail Call Optimization

Cython

```
cdef int _recurse(int n, int
count):
    if n <= 0:
        return count
    return _recurse(n-1, count+1)
```

```
def recurse(int n):
    return _recurse(n, 0)
```

timeit recurse(100)

10000000 loops, best of 3: **134 ns**
per loop

Python

```
def recurse(n, count=0):
    if n <= 0:
        return count
    return recurse(n-1, count+1)
```

timeit recurse(100)

10000 loops, best of 3: **31.3 us** per
loop

Extension Classes

```
cdef class Shrubbery:
    cdef int width, height
    def __init__(self, w, h):
        self.width = w
        self.height = h
    def describe(self):
        print "This shrubbery is", self.width, \
        "by", self.height, "cubits."
```

- Pretty simple, it's just like defining an extension class in C
- Attribute access is just a struct lookup
- Special `__cinit__` and `__dealloc__` functions

Extern C headers

```
cdef extern from "zmq.h" nogil:
    # blackbox def for zmq_msg_t
    ctypedef void * zmq_msg_t "zmq_msg_t"
...
    enum: ZMQ_PUB # 1
...
    # send/recv
    int zmq_sendmsg (void *s, zmq_msg_t *msg, int flags)
    int zmq_recvmsg (void *s, zmq_msg_t *msg, int flags)
...
    ctypedef struct zmq_pollitem_t:
        void *socket
        int fd
        short events
        short revents
```

Quick example

Cython

```
cdef extern from "math.h":  
    double sin(double )  
cpdef double cy_sin(double x):  
    return sin(x)
```

timeit cy_sin(3.14)

10000000 loops, best of 3: **116** ns
per loop

Python

```
from math import sin
```

timeit sin(3.14)

1000000 loops, best of 3: **154** ns
per loop

Quick example 2

```
cdef extern from "math.h":  
    double sin(double)
```

- This is only callable from other cython code (since it's a cdef)
- We need to tell the build step to use libm, in order to link the symbol

```
def cy_sin(double x):  
    return sin(x)
```

- This wraps the cdef call, in a python def

```
cpdef double cy_sin2(double x):  
    return sin(x)
```

- cy_sin2 is callable efficiently from cython, and from python

Wrapping a C library

- Either wrap them in defs or extension classes

```
cdef extern from "zmq.h" nogil:
```

```
    void *zmq_socket (void *context, int type)
```

```
    int zmq_sendmsg (void *socket, zmq_msg_t  
    *msg, int flags)
```

Becomes:

ZMQ Socket

```
cdef class Socket:
    cdef void *handle
    def __cinit__(self, Context context, int socket_type):
        cdef Py_ssize_t c_handle

        c_handle = context._handle
        self.socket_type = socket_type
        self.handle = zmq_socket(<void *>c_handle, socket_type)

    cpdef object send(self, object data, int flags=0):
        cdef zmq_msg_t msg

        copy_data_to_msg(data, &msg)
        rc = zmq_sendmsg(handle, &data, flags)

        if rc < 0: raise ZMQError()
```


For in range is special

Python

```
cdef int i
for i in range(100):
    pass
```

C

```
CYTHON_UNUSED int __pyx_v_i;
int __pyx_t_1;
for (__pyx_t_1 = 0; __pyx_t_1
    < 100; __pyx_t_1+=1) {
    __pyx_v_i = __pyx_t_1;
}
```

Normal for loop

Cython

```
l = list(range(100))
for i in l:
    pass
```

C

```
__pyx_t_1 = ((PyObject *)__pyx_v_1);
__Pyx_INCREF(__pyx_t_1);
__pyx_t_3 = 0;
for (;;) {
    if (__pyx_t_3 >=
        PyList_GET_SIZE(__pyx_t_1))
        break;
    __pyx_t_2 = PyList_GET_ITEM(__pyx_t_1,
__pyx_t_3);
    __Pyx_INCREF(__pyx_t_2);
    __pyx_t_3++;
    __Pyx_XDECREF(__pyx_v_i);
    __pyx_v_i = __pyx_t_2;
    __pyx_t_2 = 0;
}
```

Using Cython interface files

- Cython files come in 3 flavors .pyx, .pxd, and .pxi
- .pyx is where most of your implementation will go
- .pxd is similar to .h files
- For example pyzmq was written almost entirely in Cython, and they export a cython interface alongside a python one

Pyzmq file structure

- Pxd files along with binaries

```
paryani at eotdatabal811 in ~$ ls pyzmq/core
constants.so    __init__.pyc  _poll.so      stopwatch.so
context.pxd    libzmq.pxd   pysocket.py   version.py
context.so     message.pxd  pysocket.pyc  version.pyc
device.so     message.so   socket.pxd    _version.so
error.so      poll.py      socket.so
__init__.py   poll.pyc    stopwatch.pxd
```

Pyzmq

- libzmq.pxd contains a raw wrapping of the zmq C api
- context.pxd and socket.pxd contain definitions of pyzmq's Extension classes, Context and Socket
- These pxd files allow us to import the extension classes in Cython, and access cython specific fields and methods

Cimport

- Usable only from cython
- Very similar to import, but will only import .pxd files in the python path
- Usable like so:

```
from zmq.core.socket cimport Socket
```

Socket.pxd

```
from context cimport Context
```

```
cdef class Socket:
```

```
    """A ØMQ socket."""
```

```
    cdef void *handle # The C handle for the underlying zmq object.
```

```
    cdef public int socket_type # The ØMQ socket type - REQ, REP, etc.
```

```
    cdef public Context context # The zmq Context object that owns this.
```

```
    cdef public bint _closed # bool property for a closed socket.
```

```
    cdef dict _attrs # dict needed for *non-sockopt* get/setattr in  
    subclasses
```

```
    cdef int _pid # the pid of the process which created me (for fork  
    safety)
```

```
# cpdef methods for direct-cython access:
```

```
cpdef object send(self, object data, int flags=*, copy=*, track=*)
```

```
cpdef object recv(self, int flags=*, copy=*, track=*)
```

Simple zmq publisher

```
import zmq
```

```
context = zmq.Context()
```

```
publisher = context.socket (zmq.PUB)
```

```
publisher.bind ("tcp://*:48080")
```

```
while True:
```

```
    publisher.send('1')
```


Zmq Python client

```
import zmq
```

```
context = zmq.Context()
subscriber = context.socket (zmq.SUB)
subscriber.connect("tcp://localhost:48080")
subscriber.setsockopt (zmq.SUBSCRIBE, "")
```

```
def main_loop(subscriber):
    count = 0
    while True:
        if count > 1000000: break
        count += int(subscriber.recv())
main_loop(subscriber)
```

C equivalent

```
#include "zmq.h"
```

```
int main(int argc, char ** argv) {  
    void *context = zmq_init (1);
```

```
    // Socket to talk to clients
```

```
    void *subscriber = zmq_socket (context, ZMQ_SUB);  
    zmq_connect (subscriber, "tcp://localhost:48080");  
    zmq_setsockopt (subscriber, ZMQ_SUBSCRIBE, "", 0);
```

```
    int count = 0;
```

```
    while(1)
```

```
{
```

```
    if(count > 1000000) break;
```

```
    zmq_msg_t msg;
```

```
    zmq_msg_init (&msg);
```

```
    zmq_recv (subscriber, &msg, 0);
```

```
    count += atoi(zmq_msg_data(&msg));
```

```
    zmq_msg_close (&msg);
```

```
}
```

```
}
```

Zmq Cython Client

```
# zmq_cython_client.pyx
from zmq.core.socket cimport Socket
from zmq.core.libzmq cimport zmq_msg_t, zmq_msg_init,
    zmq_msg_close, zmq_msg_data, zmq_recvmmsg
from libc.stdlib cimport atoi

def main_loop(Socket subscriber):
    cdef int count = 0
    cdef zmq_msg_t msg
    while True:
        if count > 1000000: break
        zmq_msg_init(&msg)
        zmq_recvmmsg(subscriber.handle, &msg, 0)
        count += atoi(<char *> zmq_msg_data(&msg))
        zmq_msg_close(&msg)
```

Zmq Python client with Cython

```
import zmq
```

```
context = zmq.Context()
```

```
subscriber = context.socket (zmq.SUB)
```

```
subscriber.connect("tcp://localhost:48080")
```

```
subscriber.setsockopt (zmq.SUBSCRIBE, "")
```

```
import zmq_cython_client
```

```
zmq_cython_client.main_loop(subscriber)
```

Speed comparison

- C ~.33 second per run
- Cython ~.35 second per run
- Python ~1.8 second per run

Building

- Either done through an extension module in a setup.py:

```
from distutils.core import setup
from distutils.extension import Extension
from Cython.Distutils import build_ext
```

```
ext_modules=[
    Extension("cython_test", ["cython_test.pyx"])
]
setup(
    name = "cython_test",
    cmdclass = {"build_ext": build_ext},
    ext_modules = ext_modules
)
```

Building 2

- Or through pyximport:

```
import pyximport  
pyximport.install()
```

```
import cython_test
```

- Read the docs on it, this can get complicated
<http://docs.cython.org/src/quickstart/build.html>

Further Features in Cython

- (Relatively) easy to use GIL management for multithreading
- Also integrates nicely with numpy and the buffer interface
- Fused types (Similar to type templates)
- C++ linking and syntax



Questions?